

Machine Learning

Pham Vo

2/13/2022

Clustering with kmeans and hclust()

We will begin by making up some data to cluster. shortcut to insert code chunk: option+command+i

```
tmp <- c(rnorm(30,3), rnorm(30,-3))
x <- cbind(x=tmp,y=rev(tmp))
x
```

```
##           x           y
## [1,]  4.2347574 -4.9365848
## [2,]  3.3596015 -3.0430727
## [3,]  4.3189254 -1.8155747
## [4,]  2.3577094 -3.9550889
## [5,]  3.4140226 -3.1392119
## [6,]  3.4825113 -3.4872698
## [7,]  3.6417728 -2.1300113
## [8,]  2.8425338 -3.5283979
## [9,]  3.9547135 -2.5499144
## [10,] 2.6801855 -3.4023781
## [11,] 3.2225554 -3.7673758
## [12,] 2.6247568 -4.4102463
## [13,] 1.4888526 -3.8175338
## [14,] 3.6314556 -3.0617546
## [15,] 1.2445519 -3.8664872
## [16,] 4.0780839 -1.2648538
## [17,] 2.9736588 -1.6253573
## [18,] 2.8033944 -3.5290531
## [19,] 3.5371391 -0.7760335
## [20,] 3.3591323 -3.6637642
## [21,] 3.7945046 -4.1831013
## [22,] 2.7439267 -3.8402006
## [23,] 4.4674454 -4.2069651
## [24,] 2.4958294 -4.6194852
## [25,] 2.6650930 -2.2922960
## [26,] 2.1556406 -3.7717725
## [27,] 3.4293102 -3.0912247
## [28,] 3.9995536 -1.0218711
## [29,] 2.3447089 -0.4673843
## [30,] 1.8751953 -1.5395258
## [31,] -1.5395258  1.8751953
```

Run kmeans()

```
## K-means clustering with 2 clusters of sizes 30, 30
##
## Cluster means:
##      x          y
## 1 -3.026793   3.107384
## 2  3.107384  -3.026793
##
## Clustering vector:
## [1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
## [39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
##
## Within cluster sum of squares by cluster:
## [1] 61.71633 61.71633
## (between_SS / total_SS =  90.1 %)
##
## Available components:
```

```
## [1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
## [6] "betweenss"    "size"         "iter"         "ifault"
```

Q. what size is each cluster?

k\$size

```
## [1] 30 30
```

Q. Cluster centers

k\$centers

```
##           x           y
## 1 -3.026793  3.107384
## 2  3.107384 -3.026793
```


Q. Membership vector

```
k$cluster
```

```
## [1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1 1
## [39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

Plot our data with the clustering result

```
plot(x, col=k$cluster)
points(k$centers, col="blue", pch=16, cex=1.5)
points(x[42,1], x[42,2], col="darkgreen", pch = 16)
points(x[19,1], x[19,2], col="brown", pch = 16)
```



hclust()


```
hc <- hclust(dist(x))
hc
```

```
##
## Call:
## hclust(d = dist(x))
##
## Cluster method   : complete
## Distance         : euclidean
## Number of objects: 60
```

There is a cool and useful plot method for hclust()

```
plot(hc)
```

Cluster Dendrogram



Principal Component Analysis

Data import

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names=1)
x
```

##	England	Wales	Scotland	N.Ireland
## Cheese	105	103	103	66
## Carcass_meat	245	227	242	267
## Other_meat	685	803	750	586
## Fish	147	160	122	93
## Fats_and_oils	193	235	184	209
## Sugars	156	175	147	139
## Fresh_potatoes	720	874	566	1033
## Fresh_Veg	253	265	171	143
## Other_Veg	488	570	418	355
## Processed_potatoes	198	203	220	187
## Processed_Veg	360	365	337	334
## Fresh_fruit	1102	1137	957	674
## Cereals	1472	1582	1462	1494
## Beverages	57	73	53	47
## Soft_drinks	1374	1256	1572	1506

```
## Alcoholic_drinks      375   475   458   135
## Confectionery         54    64    62    41
```

#Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
dim(x)
```

```
## [1] 17  4
```

```
ncol(x)
```


```
## [1] 4
```

```
nrow(x)
```

```
## [1] 17
```


Let's make some typical plots

```
mycols <- rainbow(nrow(x))
barplot(as.matrix(x), col=mycols)
```



Make it side by side i.e. not stacked

```
barplot(as.matrix(x), col=mycols, beside=TRUE)
```

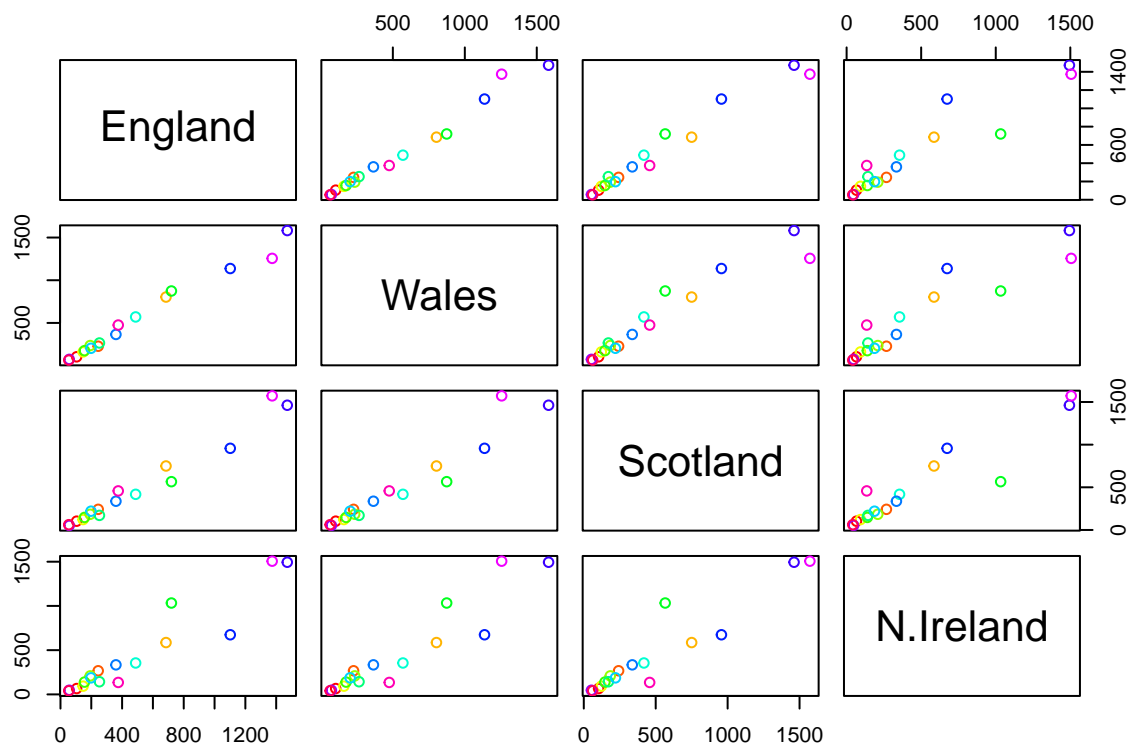


One plot that is helpful here.

#Q5: Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

No difference in values between x-axis and y-axis.

```
pairs(x, col=mycols)
```



#Q6. What is the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

PCA to rescue


Do PCR of this 17D food data. The main function in base R is called “`prcomp()`”. This function requires the transpose of our data in this case...

```
pca <- prcomp(t(x))
summary(pca)
```

```
## Importance of components:
##              PC1      PC2      PC3      PC4
## Standard deviation  324.1502 212.7478 73.87622 4.189e-14
## Proportion of Variance 0.6744 0.2905 0.03503 0.000e+00
## Cumulative Proportion 0.6744 0.9650 1.00000 1.000e+00
```

The `prcomp()` function returns a list object.

```
plot(pca)
```

The PCA plot a.k.a a PCA score plot is a plot of PC1 vs PV2. Basically using the new PCA axis to view our data.

```
attributes(pca)
```


```
## $names
## [1] "sdev"      "rotation" "center"    "scale"     "x"
##
## $class
## [1] "prcomp"
```

We will focus on “pca\$x” for this plot

```
pca$x
```

```
##           PC1          PC2          PC3          PC4
## England  -144.99315    2.532999 -105.768945  2.842865e-14
## Wales    -240.52915   224.646925   56.475555  7.804382e-13
## Scotland  -91.86934  -286.081786   44.415495 -9.614462e-13
## N.Ireland  477.39164    58.901862    4.877895  1.448078e-13
```

```
plot(pca$x[,1], pca$x[,2], xlab= "PC1", ylab="PC2")
text(pca$x[,1], pca$x[,2], labels = colnames(x))
```



```
## PCA of a RNA-Seq
```

```
url2 <- "https://tinyurl.com/expression-CSV"
rna.data <- read.csv(url2, row.names=1)
head(rna.data)
```

```
##      wt1 wt2  wt3  wt4 wt5 ko1 ko2 ko3 ko4 ko5
## gene1 439 458  408  429 420  90  88  86  90  93
## gene2 219 200  204  210 187 427 423 434 433 426
## gene3 1006 989 1030 1017 973 252 237 238 226 210
## gene4  783 792  829  856 760 849 856 835 885 894
## gene5  181 249  204  244 225 277 305 272 270 279
## gene6  460 502  491  491 493 612 594 577 618 638
```

```
pca <- prcomp(t(rna.data))
summary(pca)
```

```
## Importance of components:
##              PC1      PC2      PC3      PC4      PC5      PC6
## Standard deviation    2214.2633  88.9209  84.33908  77.74094  69.66341  67.78516
## Proportion of Variance    0.9917  0.0016  0.00144  0.00122  0.00098  0.00093
## Cumulative Proportion    0.9917  0.9933  0.99471  0.99593  0.99691  0.99784
##              PC7      PC8      PC9      PC10
## Standard deviation    65.29428  59.90981  53.20803  3.142e-13
## Proportion of Variance    0.00086  0.00073  0.00057  0.000e+00
## Cumulative Proportion    0.99870  0.99943  1.00000  1.000e+00
```

Make our PCA score plot

```
plot(pca$x[,1:2])  
text(pca$x[,1:2], labels = colnames(rna.data))
```

