

RNA-seq

Pham Vo

2/27/2022

Install Bioconductor

Import countData and colData

```
counts <- read.csv("airway_scaledcounts.csv", row.names=1)
metadata <- read.csv("airway_metadata.csv")
head(counts)
```

```
##           SRR1039508 SRR1039509 SRR1039512 SRR1039513 SRR1039516
## ENSG00000000003      723        486        904        445        1170
## ENSG00000000005         0         0         0         0         0
## ENSG00000000419      467        523        616        371        582
## ENSG00000000457      347        258        364        237        318
## ENSG00000000460       96         81         73         66        118
## ENSG00000000938         0         0         1         0         2
##           SRR1039517 SRR1039520 SRR1039521
## ENSG00000000003      1097        806        604
## ENSG00000000005         0         0         0
## ENSG00000000419      781        417        509
## ENSG00000000457      447        330        324
## ENSG00000000460       94        102        74
## ENSG00000000938         0         0         0
```

metadata

```
##           id      dex celltype      geo_id
## 1 SRR1039508 control   N61311 GSM1275862
## 2 SRR1039509 treated   N61311 GSM1275863
## 3 SRR1039512 control   N052611 GSM1275866
## 4 SRR1039513 treated   N052611 GSM1275867
## 5 SRR1039516 control   N080611 GSM1275870
## 6 SRR1039517 treated   N080611 GSM1275871
## 7 SRR1039520 control   N061011 GSM1275874
## 8 SRR1039521 treated   N061011 GSM1275875
```

How many genes are in this dataset?

```
nrow(counts)
```

```
## [1] 38694
```

Q2. How many 'control' cell lines do we have?

```
table(metadata$dex)
```

```
##  
## control treated  
##      4      4
```

```
#Extract IDs from control samples  
inds <- metadata$dex == "control"  
control.metadata <- metadata[inds,]  
control.metadata$id
```

```
## [1] "SRR1039508" "SRR1039512" "SRR1039516" "SRR1039520"
```

```
#Check columns in two tables to be the same
```

```
metadata$id == colnames(counts)
```

```
## [1] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
```

Toy differential gene expression

Note: this analysis is for demonstration only. NEVER do differential expression analysis this way!

```
# Note that the control samples are SRR1039508, SRR1039512, SRR1039516, and SRR1039520. This bit of code
```

```
control <- metadata[metadata[, "dex"]=="control",]  
control.counts <- counts[, control$id]  
control.mean <- rowSums( control.counts )/4  
head(control.mean)
```

```
## ENSG00000000003 ENSG00000000005 ENSG00000000419 ENSG00000000457 ENSG00000000460  
##      900.75      0.00      520.50      339.75      97.25  
## ENSG00000000938  
##      0.75
```

Q3. How would you make the above code in either approach more robust?

Q4. Follow the same procedure for the treated samples (i.e. calculate the mean per gene across drug treated samples and assign to a labeled vector called treated.mean)


```
treated <- metadata[metadata[, "dex"]=="treated",]
treated.mean <- rowSums( counts[ ,treated$id] )/4
names(treated.mean) <- counts$ensgene

# combine our meancount data for bookkeeping purposes
meancounts <- data.frame(control.mean, treated.mean)
colSums(meancounts)
```

```
## control.mean treated.mean
##      23005324      22196524
```

Q5 (a). Create a scatter plot showing the mean of the treated samples against the mean of the control samples. Your plot should look something like the following.

```
library(ggplot2)
ggplot(meancounts, aes(control.mean, treated.mean)) +
  geom_point()
```




Q6. Try plotting both axes on a log scale. What is the argument to plot() that allows you to do this?

```
ggplot(meancounts, aes(control.mean, treated.mean)) +
  geom_point() +
  scale_x_continuous(trans="log2")+
  scale_y_continuous(trans="log2")
```

```
## Warning: Transformation introduced infinite values in continuous x-axis
```

```
## Warning: Transformation introduced infinite values in continuous y-axis
```



```
# calculate log2foldchange, add it to our meancounts data.frame and inspect the results either with the
```

```
meancounts$log2fc <- log2(meancounts[,green"treated.mean"]/meancounts[,green"control.mean"])  
head(meancounts)
```

```
##           control.mean treated.mean    log2fc  
## ENSG000000000003      900.75      658.00 -0.45303916  
## ENSG000000000005         0.00         0.00      NaN  
## ENSG000000000419      520.50      546.00  0.06900279  
## ENSG000000000457      339.75      316.50 -0.10226805  
## ENSG000000000460       97.25       78.75 -0.30441833  
## ENSG000000000938        0.75        0.00      -Inf
```

A common rule thumb in the field in the field is to focus initially on big changes with a cutoff log2 fc of +2 or -2.

```
# The NaN is returned when you divide by zero and try to take the log. The -Inf is returned when you try
```

```
#
```

```
zero.vals <- which(meancounts[,1:2]==0, arr.ind=TRUE)
```

```
to.rm <- unique(zero.vals[,1])
mycounts <- meancounts[-to.rm,]
head(mycounts)
```

```
##               control.mean treated.mean      log2fc
## ENSG00000000003      900.75      658.00 -0.45303916
## ENSG000000000419      520.50      546.00  0.06900279
## ENSG000000000457      339.75      316.50 -0.10226805
## ENSG000000000460       97.25       78.75 -0.30441833
## ENSG000000000971     5219.00     6687.50  0.35769358
## ENSG000000001036     2327.00     1785.75 -0.38194109
```

Q7. What is the purpose of the `arr.ind` argument in the `which()` function call above? Why would we then take the first column of the output and need to call the `unique()` function?

`arr.ind` argument: return row and column indices of the TRUE values
`unique()` function: do not count any row twice if there is a 0 value in both samples.

Q8. Using the `up.ind` vector above can you determine how many up regulated genes we have at the greater than 2 fc level?

```
sum(mycounts$log2fc > 2)
```

```
## [1] 250
```

Q9. Using the `down.ind` vector above can you determine how many down regulated genes we have at the greater than 2 fc level?

```
sum(mycounts$log2fc < (-2))
```

```
## [1] 367
```

Q10. Do you trust these results? Why or why not?

Lack of statistical analysis for differences in expression of each gene.

DESeq2 analysis

Load package

```
library(DESeq2)
```

```
## Loading required package: S4Vectors
```

```
## Loading required package: stats4
```

```

## Loading required package: BiocGenerics

##
## Attaching package: 'BiocGenerics'

## The following objects are masked from 'package:stats':
##
##     IQR, mad, sd, var, xtabs

## The following objects are masked from 'package:base':
##
##     anyDuplicated, append, as.data.frame, basename, cbind, colnames,
##     dirname, do.call, duplicated, eval, evalq, Filter, Find, get, grep,
##     grepl, intersect, is.unsorted, lapply, Map, mapply, match, mget,
##     order, paste, pmax, pmax.int, pmin, pmin.int, Position, rank,
##     rbind, Reduce, rownames, sapply, setdiff, sort, table, tapply,
##     union, unique, unsplit, which.max, which.min

##
## Attaching package: 'S4Vectors'

## The following objects are masked from 'package:base':
##
##     expand.grid, I, unname

## Loading required package: IRanges

## Loading required package: GenomicRanges

## Loading required package: GenomeInfoDb

## Loading required package: SummarizedExperiment

## Loading required package: MatrixGenerics

## Loading required package: matrixStats

##
## Attaching package: 'MatrixGenerics'

## The following objects are masked from 'package:matrixStats':
##
##     colAlls, colAnyNAs, colAnys, colAvgsPerRowSet, colCollapse,
##     colCounts, colCummaxs, colCummins, colCumprods, colCumsums,
##     colDiffs, colIQRDiffs, colIQRs, colLogSumExps, colMadDiffs,
##     colMads, colMaxs, colMeans2, colMedians, colMins, colOrderStats,
##     colProds, colQuantiles, colRanges, colRanks, colSdDiffs, colSds,
##     colSums2, colTabulates, colVarDiffs, colVars, colWeightedMads,
##     colWeightedMeans, colWeightedMedians, colWeightedSds,
##     colWeightedVars, rowAlls, rowAnyNAs, rowAnys, rowAvgsPerColSet,
##     rowCollapse, rowCounts, rowCummaxs, rowCummins, rowCumprods,

```

```
## rowCumsums, rowDiffs, rowIQRDiffs, rowIQRs, rowLogSumExps,
## rowMadDiffs, rowMads, rowMaxs, rowMeans2, rowMedians, rowMins,
## rowOrderStats, rowProds, rowQuantiles, rowRanges, rowRanks,
## rowSdDiffs, rowSds, rowSums2, rowTabulates, rowVarDiffs, rowVars,
## rowWeightedMads, rowWeightedMeans, rowWeightedMedians,
## rowWeightedSds, rowWeightedVars
```

```
## Loading required package: Biobase
```

```
## Welcome to Bioconductor
```

```
##
```

```
## Vignettes contain introductory material; view with
## 'browseVignettes()'. To cite Bioconductor, see
## 'citation("Biobase")', and for packages 'citation("pkgname")'.
```

```
##
```

```
## Attaching package: 'Biobase'
```

```
## The following object is masked from 'package:MatrixGenerics':
```

```
##
```

```
## rowMedians
```

```
## The following objects are masked from 'package:matrixStats':
```

```
##
```

```
## anyMissing, rowMedians
```

```
citation("DESeq2")
```

```
##
```

```
## Love, M.I., Huber, W., Anders, S. Moderated estimation of fold change
## and dispersion for RNA-seq data with DESeq2 Genome Biology 15(12):550
## (2014)
```

```
##
```

```
## A BibTeX entry for LaTeX users is
```

```
##
```

```
## @Article{,
## title = {Moderated estimation of fold change and dispersion for RNA-seq data with DESeq2},
## author = {Michael I. Love and Wolfgang Huber and Simon Anders},
## year = {2014},
## journal = {Genome Biology},
## doi = {10.1186/s13059-014-0550-8},
## volume = {15},
## issue = {12},
## pages = {550},
## }
```

```
dds <- DESeqDataSetFromMatrix(countData=counts,
                              colData=metadata,
                              design=~dex)
```

```
## converting counts to integer mode
```

```
## Warning in DESeqDataSet(se, design = design, ignoreRank): some variables in
## design formula are characters, converting to factors
```

```
dds
```

```
## class: DESeqDataSet
## dim: 38694 8
## metadata(1): version
## assays(1): counts
## rownames(38694): ENSG000000000003 ENSG000000000005 ... ENSG00000283120
## ENSG00000283123
## rowData names(0):
## colnames(8): SRR1039508 SRR1039509 ... SRR1039520 SRR1039521
## colData names(4): id dex celltype geo_id
```

```
dds <- DESeq(dds)
```

```
## estimating size factors
```

```
## estimating dispersions
```

```
## gene-wise dispersion estimates
```

```
## mean-dispersion relationship
```

```
## final dispersion estimates
```

```
## fitting model and testing
```

```
# To get our results out of this object we can use the DESeq2 function "results()"
```

```
res <- results(dds)
res
```

```
## log2 fold change (MLE): dex treated vs control
## Wald test p-value: dex treated vs control
## DataFrame with 38694 rows and 6 columns
##      baseMean log2FoldChange      lfcSE      stat      pvalue
##      <numeric>      <numeric> <numeric> <numeric> <numeric>
## ENSG000000000003  747.1942    -0.3507030  0.168246 -2.084470  0.0371175
## ENSG000000000005    0.0000         NA         NA         NA         NA
## ENSG000000000419  520.1342     0.2061078  0.101059  2.039475  0.0414026
## ENSG000000000457  322.6648     0.0245269  0.145145  0.168982  0.8658106
## ENSG000000000460   87.6826    -0.1471420  0.257007 -0.572521  0.5669691
## ...           ...           ...           ...           ...
## ENSG00000283115   0.000000         NA         NA         NA         NA
## ENSG00000283116   0.000000         NA         NA         NA         NA
## ENSG00000283119   0.000000         NA         NA         NA         NA
## ENSG00000283120   0.974916    -0.668258   1.69456 -0.394354  0.693319
## ENSG00000283123   0.000000         NA         NA         NA         NA
##      padj
```



```
##                               <numeric>
## ENSG000000000003  0.163035
## ENSG000000000005      NA
## ENSG000000000419  0.176032
## ENSG000000000457  0.961694
## ENSG000000000460  0.815849
## ...                ...
## ENSG00000283115      NA
## ENSG00000283116      NA
## ENSG00000283119      NA
## ENSG00000283120      NA
## ENSG00000283123      NA
```

Adding annotation data

```
library("AnnotationDbi")
library("org.Hs.eg.db")
```

```
##
```

```
columns(org.Hs.eg.db)
```

```
## [1] "ACCNUM"      "ALIAS"        "ENSEMBL"      "ENSEMBLPROT"  "ENSEMBLTRANS"
## [6] "ENTREZID"    "ENZYME"       "EVIDENCE"     "EVIDENCEALL"  "GENENAME"
## [11] "GENETYPE"    "GO"           "GOALL"        "IPI"           "MAP"
## [16] "OMIM"        "ONTOLOGY"     "ONTOLOGYALL"  "PATH"          "PFAM"
## [21] "PMID"        "PROSITE"      "REFSEQ"       "SYMBOL"        "UCSCKG"
## [26] "UNIPROT"
```

```
# We will use the main mapIds function to add different identifiers to our results
res$symbol <- mapIds(org.Hs.eg.db,
  keys=row.names(res), # Our genenames
  keytype="ENSEMBL",    # The format of our genenames
  column="SYMBOL",      # The new format we want to add
  multiVals="first")
```

```
## 'select()' returned 1:many mapping between keys and columns
```

```
head(res)
```

```
## log2 fold change (MLE): dex treated vs control
## Wald test p-value: dex treated vs control
## DataFrame with 6 rows and 7 columns
##           baseMean log2FoldChange    lfcSE      stat    pvalue
##           <numeric>      <numeric> <numeric> <numeric> <numeric>
## ENSG000000000003  747.194195    -0.350703  0.168246 -2.084470  0.0371175
## ENSG000000000005    0.000000         NA         NA         NA         NA
## ENSG000000000419  520.134160     0.2061078  0.101059  2.039475  0.0414026
## ENSG000000000457  322.664844     0.0245269  0.145145  0.168982  0.8658106
```

```
## ENSG000000000460 87.682625 -0.1471420 0.257007 -0.572521 0.5669691
## ENSG000000000938 0.319167 -1.7322890 3.493601 -0.495846 0.6200029
##                padj      symbol
##                <numeric> <character>
## ENSG000000000003 0.163035      TSPAN6
## ENSG000000000005      NA      TNMD
## ENSG000000000419 0.176032      DPM1
## ENSG000000000457 0.961694      SCYL3
## ENSG000000000460 0.815849      C1orf112
## ENSG000000000938      NA      FGR
```

Q11. Run the `mapIds()` function two more times to add the Entrez ID and UniProt accession and GENENAME as new columns called `res$entrez`, `res$uniprot` and `res$genename`.

```
res$entrez <- mapIds(org.Hs.eg.db,
                    keys=row.names(res),
                    column="ENTREZID",
                    keytype="ENSEMBL",
                    multiVals="first")
```

```
## 'select()' returned 1:many mapping between keys and columns
```

```
res$uniprot <- mapIds(org.Hs.eg.db,
                    keys=row.names(res),
                    column="UNIPROT",
                    keytype="ENSEMBL",
                    multiVals="first")
```

```
## 'select()' returned 1:many mapping between keys and columns
```

```
res$genename <- mapIds(org.Hs.eg.db,
                    keys=row.names(res),
                    column="GENENAME",
                    keytype="ENSEMBL",
                    multiVals="first")
```

```
## 'select()' returned 1:many mapping between keys and columns
```

```
head(res)
```

```
## log2 fold change (MLE): dex treated vs control
## Wald test p-value: dex treated vs control
## DataFrame with 6 rows and 10 columns
##                baseMean log2FoldChange      lfcSE      stat      pvalue
##                <numeric>    <numeric> <numeric> <numeric> <numeric>
## ENSG000000000003 747.194195    -0.3507030 0.168246 -2.084470 0.0371175
## ENSG000000000005 0.000000      NA      NA      NA      NA
## ENSG000000000419 520.134160     0.2061078 0.101059 2.039475 0.0414026
## ENSG000000000457 322.664844     0.0245269 0.145145 0.168982 0.8658106
## ENSG000000000460 87.682625    -0.1471420 0.257007 -0.572521 0.5669691
## ENSG000000000938 0.319167    -1.7322890 3.493601 -0.495846 0.6200029
```

```
##           padj      symbol      entrez      uniprot
##      <numeric> <character> <character> <character>
## ENSG00000000003  0.163035      TSPAN6       7105  AOA024RCI0
## ENSG00000000005      NA      TNMD        64102  Q9H2S6
## ENSG000000000419  0.176032      DPM1        8813  060762
## ENSG000000000457  0.961694      SCYL3       57147  Q8IZE3
## ENSG000000000460  0.815849  C1orf112     55732  AOA024R922
## ENSG000000000938      NA      FGR         2268  P09769
##
##           genename
##      <character>
## ENSG00000000003      tetraspanin 6
## ENSG00000000005      tenomodulin
## ENSG000000000419 dolichyl-phosphate m..
## ENSG000000000457 SCY1 like pseudokina..
## ENSG000000000460 chromosome 1 open re..
## ENSG000000000938 FGR proto-oncogene, ..
```

```
# Arrange and view the results by the adjusted p-value
```

```
ord <- order( res$padj )
#View(res[ord,])
head(res[ord,])
```

```
## log2 fold change (MLE): dex treated vs control
## Wald test p-value: dex treated vs control
## DataFrame with 6 rows and 10 columns
##           baseMean log2FoldChange      lfcSE      stat      pvalue
##      <numeric>      <numeric> <numeric> <numeric> <numeric>
## ENSG00000152583   954.771      4.36836  0.2371268   18.4220  8.74490e-76
## ENSG00000179094   743.253      2.86389  0.1755693   16.3120  8.10784e-60
## ENSG00000116584  2277.913     -1.03470  0.0650984  -15.8944  6.92855e-57
## ENSG00000189221  2383.754      3.34154  0.2124058   15.7319  9.14433e-56
## ENSG00000120129  3440.704      2.96521  0.2036951   14.5571  5.26424e-48
## ENSG00000148175 13493.920      1.42717  0.1003890   14.2164  7.25128e-46
##
##           padj      symbol      entrez      uniprot
##      <numeric> <character> <character> <character>
## ENSG00000152583  1.32441e-71      SPARCL1       8404  AOA024RDE1
## ENSG00000179094  6.13966e-56      PER1        5187  015534
## ENSG00000116584  3.49776e-53      ARHGEF2      9181  Q92974
## ENSG00000189221  3.46227e-52      MAOA        4128  P21397
## ENSG00000120129  1.59454e-44      DUSP1       1843  B4DU40
## ENSG00000148175  1.83034e-42      STOM        2040  F8VSL7
##
##           genename
##      <character>
## ENSG00000152583      SPARC like 1
## ENSG00000179094 period circadian reg..
## ENSG00000116584 Rho/Rac guanine nucl..
## ENSG00000189221 monoamine oxidase A
## ENSG00000120129 dual specificity pho..
## ENSG00000148175      stomatin
```

```
# Write out the ordered significant results with annotations
```

```
write.csv(res[ord,], "deseq_results.csv")
```


Summary figure: Volcano plot

```
# Setup our custom point color vector
mycols <- rep("gray", nrow(res))
mycols[ abs(res$log2FoldChange) > 2 ] <- "red"

inds <- (res$padj < 0.01) & (abs(res$log2FoldChange) > 2 )
mycols[ inds ] <- "blue"

# Volcano plot with custom colors
plot( res$log2FoldChange, -log(res$padj),
      col=mycols, ylab="-Log(P-value)", xlab="Log2(FoldChange)" )

# Cut-off lines
abline(v=c(-2,2), col="gray", lty=2)
abline(h=-log(0.1), col="gray", lty=2)
```



Pathway analysis

```
library(pathview)
```

```
## #####
```

```
## Pathview is an open source software package distributed under GNU General
## Public License version 3 (GPLv3). Details of GPLv3 is available at
## http://www.gnu.org/licenses/gpl-3.0.html. Particullary, users are required to
## formally cite the original Pathview paper (not just mention it) in publications
## or products. For details, do citation("pathview") within R.
##
## The pathview downloads and uses KEGG data. Non-academic uses may require a KEGG
## license agreement (details at http://www.kegg.jp/kegg/legal.html).
## #####
```

```
library(gage)
```

```
##
```

```
library(gageData)
```

```
data(kegg.sets.hs)
```

```
# Examine the first 2 pathways in this kegg set for humans
head(kegg.sets.hs, 2)
```

```
## $'hsa00232 Caffeine metabolism'
## [1] "10" "1544" "1548" "1549" "1553" "7498" "9"
##
## $'hsa00983 Drug metabolism - other enzymes'
## [1] "10" "1066" "10720" "10941" "151531" "1548" "1549" "1551"
## [9] "1553" "1576" "1577" "1806" "1807" "1890" "221223" "2990"
## [17] "3251" "3614" "3615" "3704" "51733" "54490" "54575" "54576"
## [25] "54577" "54578" "54579" "54600" "54657" "54658" "54659" "54963"
## [33] "574537" "64816" "7083" "7084" "7172" "7363" "7364" "7365"
## [41] "7366" "7367" "7371" "7372" "7378" "7498" "79799" "83549"
## [49] "8824" "8833" "9" "978"
```

```
foldchanges = res$log2FoldChange
names(foldchanges) = res$entrez
head(foldchanges)
```

```
##          7105          64102          8813          57147          55732          2268
## -0.35070302          NA  0.20610777  0.02452695 -0.14714205 -1.73228897
```

```
# Get the results
keggres = gage(foldchanges, gsets=kegg.sets.hs)
```

```
# Look at the object returned from gage()
attributes(keggres)
```

```
## $names
## [1] "greater" "less" "stats"
```

```
# Look at the first three down (less) pathways
head(keggres$less, 3)
```

```
##
##          p.geomean stat.mean      p.val
## hsa05332 Graft-versus-host disease 0.0004250461 -3.473346 0.0004250461
## hsa04940 Type I diabetes mellitus 0.0017820293 -3.002352 0.0017820293
## hsa05310 Asthma 0.0020045888 -3.009050 0.0020045888
##
##          q.val set.size      exp1
## hsa05332 Graft-versus-host disease 0.09053483      40 0.0004250461
## hsa04940 Type I diabetes mellitus 0.14232581      42 0.0017820293
## hsa05310 Asthma 0.14232581      29 0.0020045888
```

```
# Look at the first three up (more) pathways
head(keggres$greater, 3)
```

```
##
##          p.geomean stat.mean      p.val
## hsa00500 Starch and sucrose metabolism 0.003306262 2.772644 0.003306262
## hsa00330 Arginine and proline metabolism 0.012317455 2.280002 0.012317455
## hsa04910 Insulin signaling pathway 0.017110962 2.129511 0.017110962
##
##          q.val set.size      exp1
## hsa00500 Starch and sucrose metabolism 0.7042337      52 0.003306262
## hsa00330 Arginine and proline metabolism 0.7774866      54 0.012317455
## hsa04910 Insulin signaling pathway 0.7774866     138 0.017110962
```

```
# Make a pathway plot with our RNA-Seq expression results shown in color
```

```
pathview(gene.data=foldchanges, pathway.id="hsa05310")
```

```
## 'select()' returned 1:1 mapping between keys and columns
```

```
## Info: Working in directory /Users/phamvo/Desktop/Bioinformatics/week 8/RNA-seq
```

```
## Info: Writing image file hsa05310.pathview.png
```

```
# play with the other input arguments to pathview() to change the display in various ways including gen
```

```
# A different PDF based output of the same data
```

```
pathview(gene.data=foldchanges, pathway.id="hsa05310", kegg.native=FALSE)
```

```
## 'select()' returned 1:1 mapping between keys and columns
```

```
## Info: Working in directory /Users/phamvo/Desktop/Bioinformatics/week 8/RNA-seq
```

```
## Info: Writing image file hsa05310.pathview.pdf
```